

Predictive Capacity of Multiple Linear Regression for Patient Characteristics with Hospital Length of
Stay

Joanne Senoren

Master of Science, Data Analytics

Table of Contents

<i>Part I: Research Question and Goals</i>	3
A1. Research Question	3
A2. Goals of Analysis.....	3
<i>Part II: Model and Programming Justifications</i>	4
B1. Four Assumptions of Multiple Linear Regression.....	4
B2. Python Justification.....	4
B3. Why Use Multiple Linear Regression Model for This Analysis	5
<i>Part III: Data Preparation</i>	5
C1. Data Cleaning Goals	5
C2. Dependent and Independent Variables Summary Statistics	8
C3. Univariate and Bivariate Visualizations	11
C4. Data Transformation (Data Wrangling).....	18
C5. CSV File of Prepared Data.....	21
<i>Part IV: Model Comparison and Analysis</i>	21
D1. Initial Model	21
D2. Linear Model Reduction Justification.....	24
D3. Reduced Model.....	26
E1. Models Comparison	27
E2. Screenshots of Outputs and Calculation.....	29
E3. Executable Code.....	30
<i>Part V: Data Summary and Implications</i>	31
F1. Results of Data Analysis	31
F2. Actionable Recommendations.....	34
<i>Part VI: Demonstration</i>	35
G. Sources for Third-Party Code	35
H. References.....	36

Part I: Research Question and Goals

A1. Research Question

Hospital length of stay (LOS) is a primary measure used to determine the quality of care in hospitals worldwide. Other factors, such as readmission and mortality rate, are often used to determine areas of improvement across administration, patient care, facilities, and spread of resources (Lingsma et al., 2018). However, we should also consider specific patient-level factors to help guide a hospital's performance and treatment abilities.

This consideration leads us to this paper's research question: *What factors regarding a patient's hospital activities, medical history, and demographic background correlate with the length of their initial hospital stay?* Finding out which factors surrounding a patient's hospital stay can help direct the course of action across patient prioritization, improved care, and streamlined resources. Patients who must spend unnecessary additional days in the hospital are more at risk for medical complications. Additionally, prolonged hospital stays put more financial stress on the hospital and healthcare system. (Tipton et al., 2021)

A2. Goals of Analysis

The overall goal of this analysis is to determine whether patient-level factors in the dataset influence a patient's hospital stay length. To reach this goal of this analysis, I will select which relevant variables in the dataset we can observe against 'Initial_days' (hospital stay) and develop a multiple linear regression model to specify which variables among those selected show a significant correlation with patient hospital stays. After generating a model, it will be scaled down to determine which combination of features can predict the length of hospital stay. Identifying these factors regarding patient demographics and medical information from the dataset can help increase efficiency in patient turnover, improve patient care, streamline hospital costs, and guide staffing decisions.

Part II: Model and Programming Justifications

B1. Four Assumptions of Multiple Linear Regression

The following are four assumptions of multiple linear regression: linearity, non-multicollinearity, homoscedasticity, and multivariate normality.

Linearity assumes a linear relationship between the response variable and the predictor variable. When selecting predictor variables, I established a loose correlation coefficient threshold for this analysis. Non-multicollinearity means that the model consists of predictor variables that are not highly correlated to each other to avoid redundancy that can affect coefficients, making predictions unreliable. Homoscedasticity demonstrates a constant variance among the residuals throughout the model, which helps us establish that the model is reliable. Multivariate normality is the assumption that the residuals in a linear regression model are normally distributed (Bobbitt, 2021). Whether the model meets these assumptions is discussed in the data analysis section of the paper.

B2. Python Justification

My programming language of choice for this linear regression model is Python within the Jupyter Notebook environment. The first rationale for Python is that it is semantically easier to read than R. I also keep my code and annotations more organized with Python. By maintaining well-structured code, a reader can quickly understand what is happening in addition to the annotations that can help guide on what could happen next. The second justification for Python is my familiarity with several Python packages I benefitted from using for building this analysis. Packages such as NumPy, pandas, matplotlib, seaborn, and statsmodels were helpful in my analysis across data cleaning, preparation, wrangling, modeling, and feature selection. I used these packages repeatedly throughout the entire analytical process. Table 1 below demonstrates how each package contributes to the analysis goals, data cleaning, wrangling, and modeling.

Table 1

Python Packages and Uses for Analysis

Package Name	Package Use in Analysis
NumPy	- used to work with arrays - mathematical functions for calculations
pandas	- create DataFrames - manipulate DataFrames - generate statistical summaries
matplotlib	- establish and customize subplots for visualizations - simple histograms and bar plots
seaborn	- customize boxplots, strip plot, heatmap, residual plot
statsmodels	- create a fitted regression model - check VIF scores - create a QQ plot

B3. Why Use Multiple Linear Regression Model for This Analysis

We can analyze a variety of continuous and categorical variables against the response variable 'Initial_days,' a continuous variable with the float data type. The float data type means that it is measured with any value within a range of any given values (Bruce et al., 2020). Multiple linear regression utilizes a specific calculation to determine if a combination of explanatory variables (continuous or categorical) can predict a continuous response variable. Since we are trying to find and select what categorical and continuous factors about a patient contribute to the continuous length of their hospital stay, multiple linear regression is an appropriate model to explore this type of analysis.

Part III: Data Preparation

C1. Data Cleaning Goals

The primary data cleaning goal for this dataset is to prepare the data for analysis by doing the following steps to eliminate data noise and gain more clarity about the dataset. Firstly, we should check for nulls and impute them if necessary (Figure 1). Second, duplicates should be addressed and removed as necessary (Figure 2). The third step would be to check the provided dictionary against the dataset to ensure the unique values in the variables are appropriate and make any necessary changes (Figure 3). Additional cleaning steps would include updating column names to match the dictionary to make the

dataset more precise (Figure 4). A more detailed look at the data also includes looking at outliers among continuous variables by generating visualizations and then individually considering each variable's outliers and whether to treat them based on realistic ranges (Figure 5). An example in Figure 6 shows how outliers for the variable 'Income' were examined.

These data preparation steps are essential tasks prior to predictive modeling, such as linear regression, because ensuring that the fitted model has accurate data is important. We gain an understanding of how reliable the model will be by addressing issues such as missing values, duplicates, and outliers. To see the annotated code used to clean the dataset in detail, please see the attached file named "D208_MultiLinearRegression.ipynb" and view the section with the header "Data Cleaning and Preparation."

Figure 1

Code for Checking for Nulls in Dataset

Checking for Nulls/NAs ¶

Confirmed in dictionary that there are no values that should be entered as 'None'

```
# Look for number of rows with minimum of one missing value
(df.isnull().sum(axis=1) > 0).sum()
0
```

Figure 2

Code for Checking Duplicates

Check for Duplicates

```
df.duplicated().value_counts()
False    10000
Name: count, dtype: int64
```

Figure 3

Using `.unique()` to Check Input Values

Compare Unique Value Samples Against Dictionary

```
# Check unique values against dictionary
for col in df:
    print(col, ':', df[col].unique())
```

Figure 4

Code to Update Specific Column Names

Specify Survey Variable Names

Change from 'Item...' to respective dictionary names

```
# Establish dictionary to map survey names
name_dict = {'Item1': 'TimelyAdmission', 'Item2': 'TimelyTreatment', 'Item3': 'TimelyVisits', 'Item4': 'Reliability', 'Item5': 'Options', 'Item6': 'TreatmentHours',
             'Item7': 'CourteousStaff', 'Item8': 'ActiveListening'}

# Rename the dataframe and change the variable names so it persists
df.rename(columns=name_dict, inplace=True)

print(df.info())
```

Figure 5

Code for Checking Outliers

Look at Outliers Among Continuous Variables

Generate box plots to identify outliers quickly and then individually examine them to determine treatment.

```
cont_var = ['Population', 'Children', 'Age', 'Income', 'VitD_levels', 'Doc_visits', 'Full_meals_eaten',
            'vitD_supp', 'Initial_days', 'TotalCharge', 'Additional_charges']

# Set figure options
fig, axes = plt.subplots(4, 3, figsize=(20, 20))
fig.subplots_adjust(hspace=.5, wspace=.25)

# For loop over list of variables to generate box plots
for var, ax in zip(cont_var, axes.flat):
    ax.set_label(var)
    sns.boxplot(data=df, x=var, ax=ax)

axes[3,2].remove() # Removes unused chart
# Show figure
plt.show()
```

Figure 6

Example Consideration for 'Income'

```

### Income Outliers
This variable has the widest range of values with minimum of 154 and maximum of 207,249.
An explanation for the minimum income could be that patient has reported unemployment benefits.
On the maximum end, it's possible for someone to earn above 200,000.
The outliers are outside the upper limit so let's count how many there are.

: # Gets Q1 and Q3 values
q1, q3 = np.percentile(df['Income'], [25, 75])

# Calculate interquartile range
iqr = q3 - q1

# Calculate upper limit
upper = q3 + (1.5 * iqr)

# Set variable for count
upper_count = 0

# Count outliers in Population
for x in df['Income']:
    if x > upper:
        upper_count += 1

print('Upper Outliers:', upper_count)

obs_perc = upper_count/len(df)
print('Percentage of dataset: {:.1%}'.format(obs_perc))
print(df['Income'].nlargest(n=upper_count))

Upper Outliers: 326
Percentage of dataset: 3.3%
8386 287249.10
841 284542.41
8598 283774.60
6468 197675.00
1778 197576.18
...
9791 186592.00
7391 186691.20
4027 106671.88
9596 106592.60
957 106521.92
Name: Income, Length: 326, dtype: float64

Outliers start at 106,521 to the maximum 207,249. These are realistic salaries and self-reported data. So we will keep Income the same.

```

C2. Dependent and Independent Variables Summary Statistics

The following table (Table 2) shows the independent variables selected for this analysis. They are variables that describe a patient's demographic background, hospital interactions, and medical history. These will be modeled against the dependent variable, 'Initial_days,' which is the length of stay (in days) of the patient's first hospitalization.

Table 2

Independent Variables Selected for MLR

Name	Variable Type	Data Type	Description
Age	Independent	Quantitative	patient's age upon admission
VitD_levels	Independent	Quantitative	patient's vitamin D levels (ng/ml)
TotalCharge	Independent	Quantitative	average per patient based on total charge and hospital stay
Doc_visits	Independent	Quantitative	number of times the doctor visited the patient during the hospital stay
Full_meals_eaten	Independent	Quantitative	number of full meals the patient ate during the hospital stay
Complication_risk	Independent	Categorical	level of patient complication risk
Initial_admin	Independent	Categorical	type of initial hospital admission
Services	Independent	Categorical	type of services the patient received while hospitalized
Soft_drink	Independent	Categorical	whether or not the patient drinks 3+ sodas a day
HighBlood	Independent	Categorical	Yes' or 'No' patient has high blood pressure

Overweight	Independent	Categorical	Yes' or 'No' patient is overweight
Arthritis	Independent	Categorical	Yes' or 'No' patient has arthritis
Diabetes	Independent	Categorical	Yes' or 'No' patient has diabetes
BackPain	Independent	Categorical	Yes' or 'No' patient has chronic back pain
Reflux_esophagitis	Independent	Categorical	Yes' or 'No' patient has reflux esophagitis
Asthma	Independent	Categorical	Yes' or 'No' patient has asthma

Figure 7 below shows the summary statistics of the selected quantitative independent variables and the dependent variable.

Figure 7

Summary Statistics for Dependent Variable and Continuous Independent Variables

<code>## Dependent Variable Summary Statistics</code>		<code>df[['VitD_levels', 'TotalCharge', 'Age', 'Doc_visits', 'Full_meals_eaten']].describe()</code>				
		<code>VitD_levels</code>	<code>TotalCharge</code>	<code>Age</code>	<code>Doc_visits</code>	<code>Full_meals_eaten</code>
<code>count</code>	10000.000000	<code>count</code>	10000.000000	10000.000000	10000.000000	10000.000000
<code>mean</code>	34.455299	<code>mean</code>	17.964262	5312.172769	53.511700	5.012200
<code>std</code>	26.309341	<code>std</code>	2.017231	2180.393838	20.638538	1.045734
<code>min</code>	1.001981	<code>min</code>	9.806483	1938.312067	18.000000	1.000000
<code>25%</code>	7.896215	<code>25%</code>	16.626439	3179.374015	36.000000	4.000000
<code>50%</code>	35.836244	<code>50%</code>	17.951122	5213.952000	53.000000	5.000000
<code>75%</code>	61.161020	<code>75%</code>	19.347963	7459.699750	71.000000	6.000000
<code>max</code>	71.981490	<code>max</code>	26.394449	9180.728000	89.000000	7.000000

The figure above shows the summary statistics of the quantitative variables utilized in the model.

The summary considers the 10,000 observations in the dataset as indicated in the 'count' row. The mean rows display the average value for each variable, the sum of all values divided by the number of rows. For example, the 'Initial_days' mean of 34.45 tells us that, on average, a patient in the dataset has had an initial hospitalization stay of about 34 days. The 'std' is the standard deviation or spread of the data around its mean. One standard deviation in 'Initial_days' is ~26 days, which tells us that a stay of 50 days is within one standard deviation. 'Min' and 'max' stand for minimum and maximum values, respectively. For 'Initial_days,' we see that the minimum value is 1 day of hospitalization, and the maximum value is about 72 days. The following univariate and bivariate sections provide additional summary statistics.

50% is our median, or middle value, about 36 days. The first and third quartiles are 25% and 75%, respectively. 50% of the data is inside the 25% and 75% markers. The other 50% are split in half outside the first and third quartiles. Since the first quartile has a value of 7.8, we can say that 25% of the data has a hospital stay of less than 7.8 days. The third quartile represents the higher values, so we can say that 25% of hospital stays are more than 61 days. The definitions apply to the rest of the numeric variables.

As the D208 PA step-by-step guide mentioned, the general statistics for categorical variables are generated in counts and percentages. Therefore, the figure below shows counts of each category per variable and their percentage within the dataset.

Figure 8

Categorical Variables Summary Statistics

```
: category_vars = ['Complication_risk', 'Initial_admin', 'Services', 'Soft_drink', 'HighBlood', 'Overweight', 'Arthritis', 'Diabetes', 'BackPain', 'Reflux_esophagitis', 'Asthma']
```

```
df[category_vars].describe()
```

	Complication_risk	Initial_admin	Services	Soft_drink	HighBlood	Overweight	Arthritis	Diabetes	BackPain	Reflux_esophagitis	Asthma
count	10000	10000	10000	10000	10000	10000	10000	10000	10000	10000	10000
unique	3	3	4	2	2	2	2	2	2	2	2
top	Medium	Emergency Admission	Blood Work	No	No	Yes	No	No	No	No	No
freq	4517	5060	5265	7425	5910	7094	6426	7262	5886	5865	7107

```
# Counts and Percentages
```

```
for col in df[category_vars]:
    counts = df[col].value_counts()
```

```
percs = df[col].value_counts(normalize=True).mul(100).round(1).astype(str) + '%'
```

```
print("\n")
```

```
print(pd.concat([counts,percs], axis=1, keys=['count', 'percentage']))
```

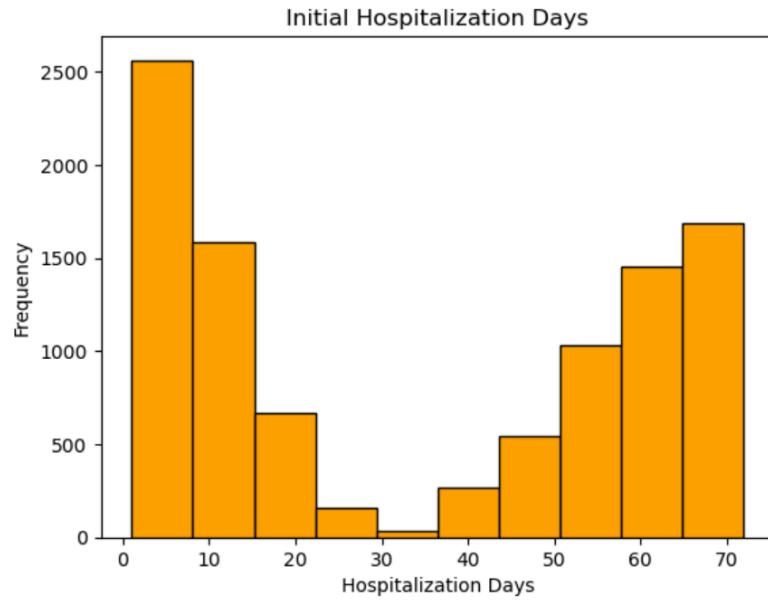
Complication_risk		count percentage		HighBlood		count percentage	
Medium		4517 45.2%		No		5910 59.1%	
High		3358 33.6%		Yes		4090 40.9%	
Low		2125 21.2%					
Initial_admin		count percentage		Overweight		count percentage	
Emergency Admission		5060 50.6%		Yes		7094 70.9%	
Elective Admission		2504 25.0%		No		2906 29.1%	
Observation Admission		2436 24.4%					
Services		count percentage		Arthritis		count percentage	
Blood Work		5265 52.6%		No		6426 64.3%	
Intravenous		3130 31.3%		Yes		3574 35.7%	
CT Scan		1225 12.2%					
MRI		380 3.8%					
Soft_drink		count percentage		Diabetes		count percentage	
No		7425 74.2%		No		7262 72.6%	
Yes		2575 25.8%		Yes		2738 27.4%	
BackPain		count percentage		Reflux_esophagitis		count percentage	
No				No		5865 58.6%	
Yes				Yes		4135 41.3%	
Asthma		count percentage					
No				No		7107 71.1%	
Yes				Yes		2893 28.9%	

The first chart in Figure 8 shows each variable's mode or most frequented category. The values in 'Soft_drink,' 'HighBlood,' 'Overweight,' 'Arthritis,' 'Diabetes,' 'BackPain,' 'Reflux_esophagitis,' and 'Asthma' are Boolean values consisting of Yes and No. It is interesting to note that the mode for these variables is mostly 'No' with an exception for Overweight, where 'Yes' is the mode. The third screenshot in Figure 8 goes over the percentages for each categorical variable in relation to the dataset. 'Initial_admin,' for example, tells us that 50.6% of hospitalized patients were admitted on an emergency admission status. The univariate section also reflects these percentages because of the limited statistics that categorical variables can provide.

C3. Univariate and Bivariate Visualizations

Figure 9

Univariate Statistic for Response Variable



The variable 'Initial_days' and its univariate statistics are shown in Figure 9 above. The generated histogram shows that the dependent variable demonstrates a bimodal distribution with two peaks. This tells us that there are two distinct groups represented in the variable. Although we do not have more details about them, we can see that there are two sets of concentrations where many patients stay less than 10 days or more than 65 days. At this point, it's important to note that something else may be the cause of two subsets within one continuous variable. The histogram also tells us that the minimum and maximum values fall within two standard deviations and that no outliers exist.

Figure 10

Univariate and Bivariate Statistics Continuous Explanatory Variables

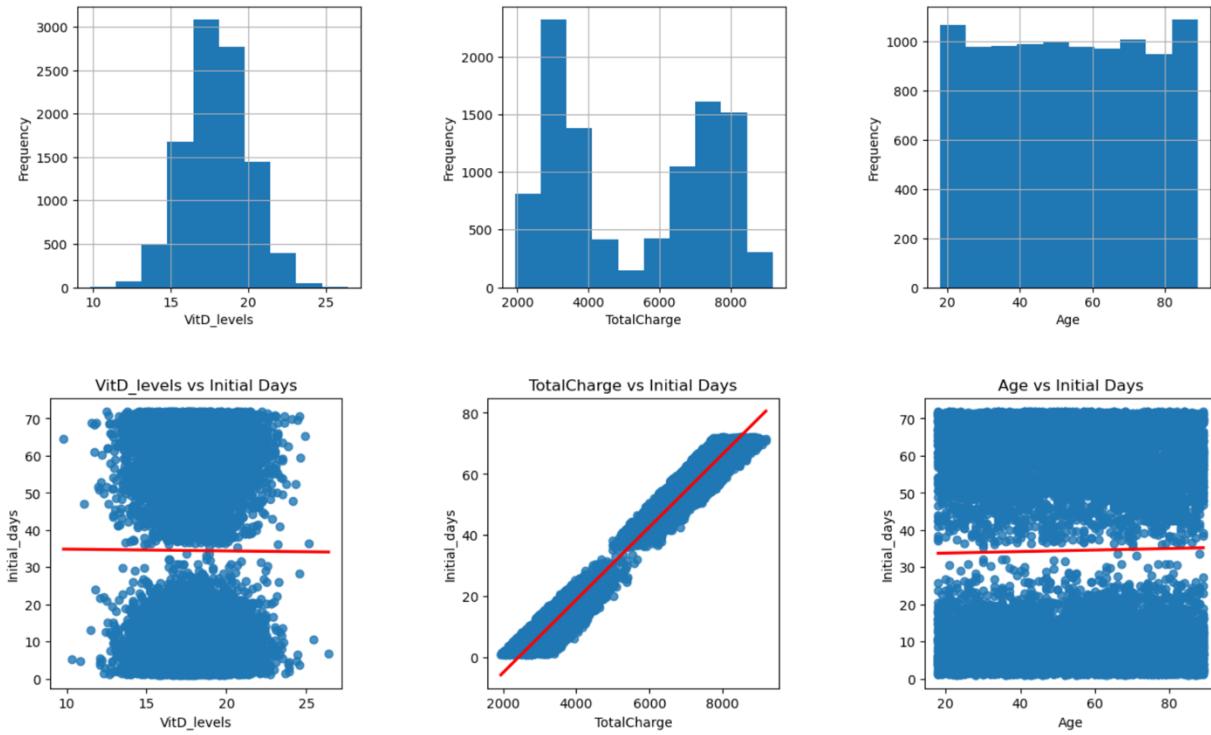
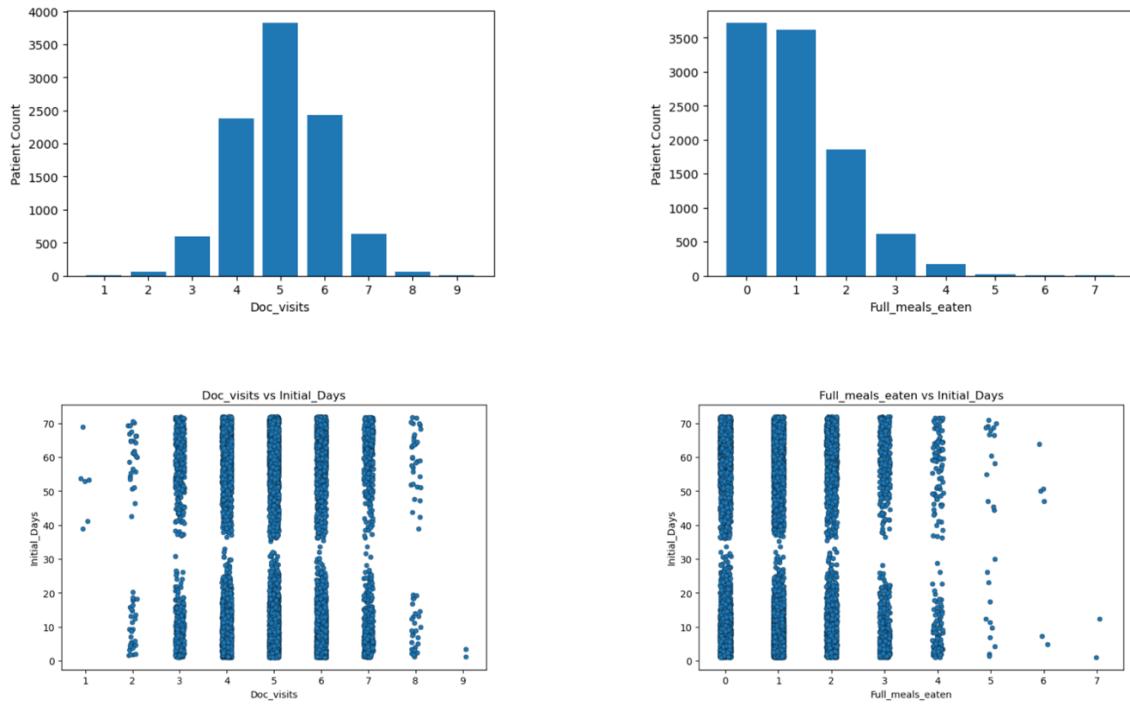


Figure 10 shows two rows of visualizations. The first row is the univariate visualizations for the continuous variables, and the second row consists of their bivariate visualizations against 'Initial_days.' The 'VitD_levels' univariate graph demonstrates a normal distribution due to the approximate bell-shaped curve of the data. 'TotalCharge' shows a bimodal distribution with two peaks on either end of the histogram. 'Age' has a uniform distribution, which tells us that the patients have nearly equal subsets of age groups.

The red line in the second row within the scatterplots is a regression line that plots the relationship between the response variable and explanatory variables. 'TotalCharge' demonstrates a potential linear relationship with 'Initial_days.' We also know that the univariate graphs for 'TotalCharge' and 'Initial_days' are bimodal, indicating their positive correlation. The other two scatterplots show a nearly flat regression line or minimum relationship.

Figure 11

Univariate and Bivariate Statistics Discrete Explanatory Variables



The first row in Figure 11 shows the univariate bar graphs for two selected discrete variables. The data in 'Doc_visits' is normally distributed, and most patients visit their doctor approximately five times during their hospitalization. 'Full_meals_eaten' is right-skewed or positively skewed, where more patients eat none, partial, or one meal daily. The bivariate strip plots on the second row show some interesting points. 'Doc_visits' against 'Initial_days' shows that most people who stay as little as 1 day or as much as 70 days are visited by a doctor 5 times. The lower number of visits demonstrates less distribution of patients. It's interesting to note that those who are only visited by a doctor once have stayed in the hospital for more than 30 days. It is also notable in the 'Full_meals_eaten' vs. 'Initial_days' strip plot that many patients, regardless of their hospital stay length, have not eaten full meals. An assumption would be that short-term patients do not get to eat full meals before discharge, and long-term patients get the time to eat more than one full meal. However, the graph doesn't support this theory. This tells me there might be something specific about the patients in the dataset.

Figure 12

Univariate Statistics Categorical Explanatory Variables

Figure 12 visualizes the summary statistics across the selected categorical variables in section C2. The percentages are showcased across the individual pie charts. We can quickly add additional details to the written summary statistics. For example, even though the summary statistic for 'HighBlood' told us that the mode is "No," we can easily see that the split between the patients who have and don't have high blood pressure is nearly fifty percent.

We can also better understand the non-Boolean categorical variables and note that across 'Complication_risk,' 'Initial_admin,' and 'Services,' the mode makes up approximately half of the patients in the data.

Figure 13

Bivariate Statistics Categorical Explanatory Variables

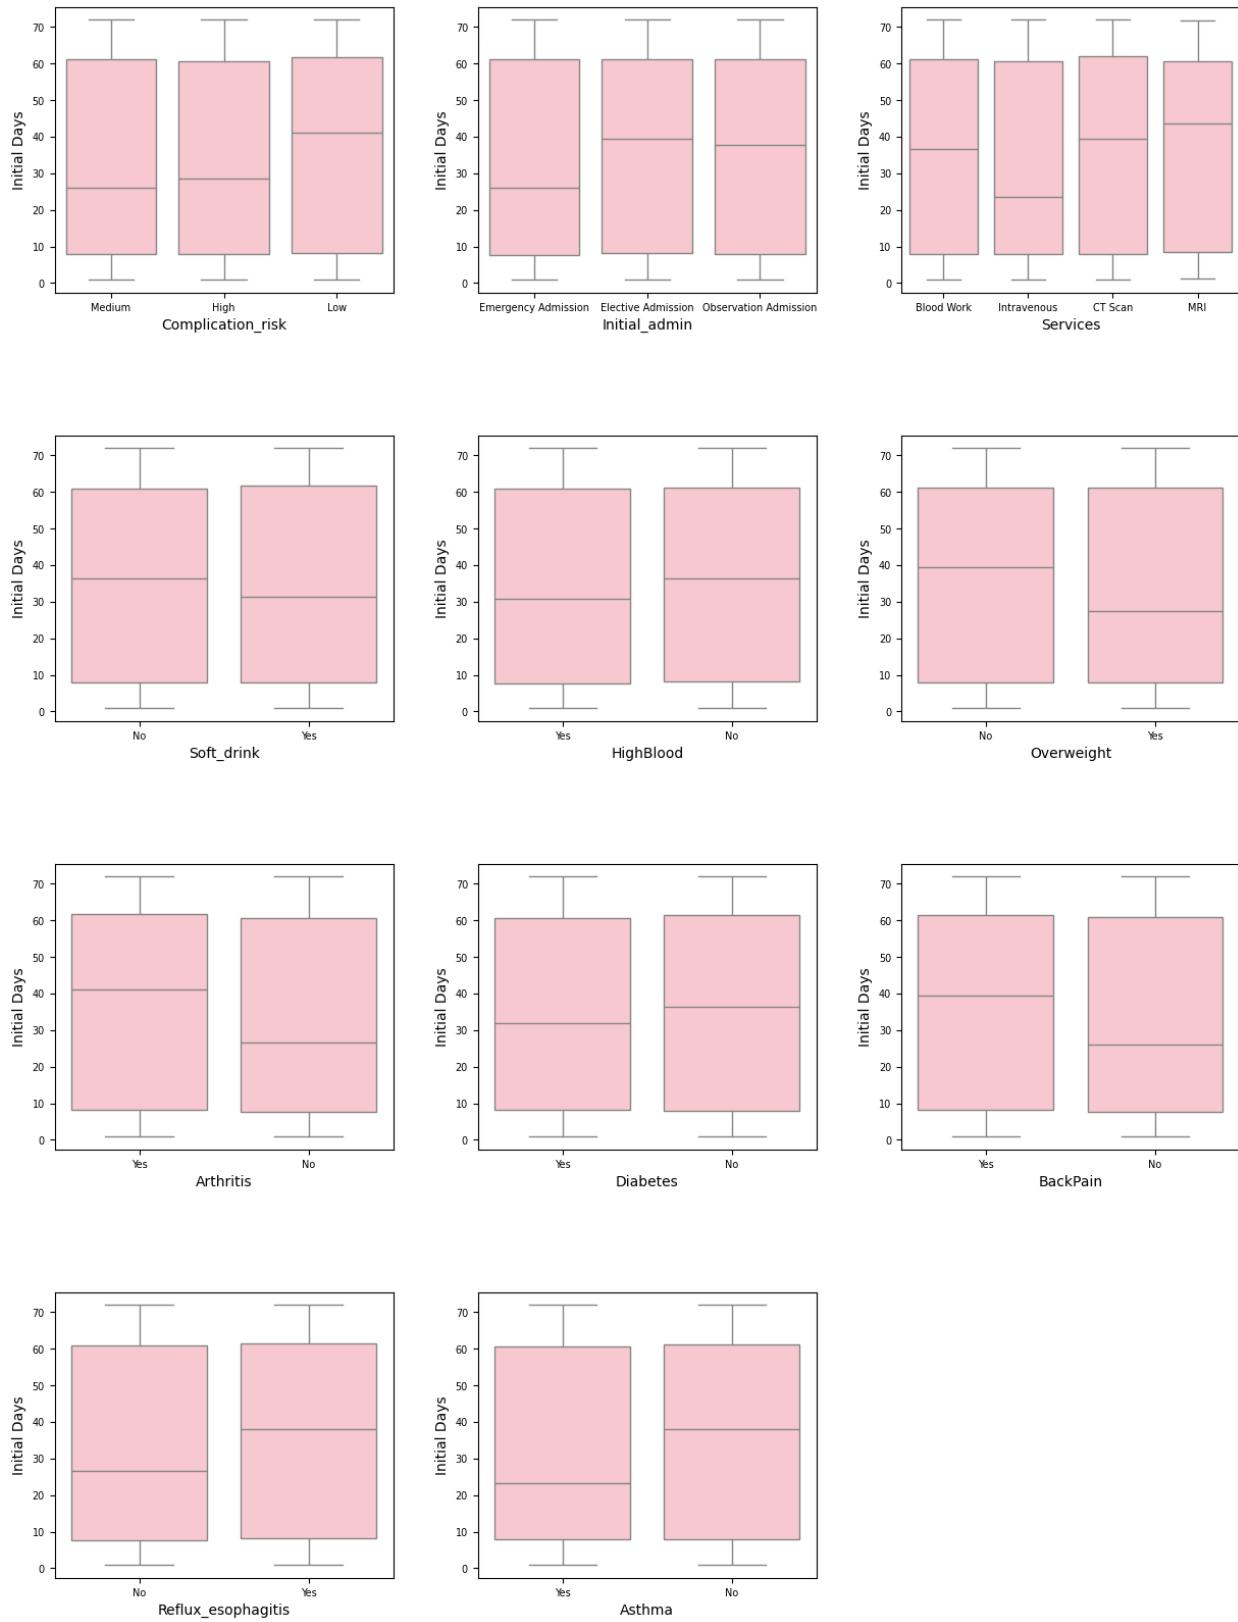


Figure 13 above shows us the bivariate visualizations against 'Initial_days.' The boxplots add even more information to the data we are exploring by demonstrating the individual medians of each group within the categorical variable. We see some notable information here concerning hospital stay and the patient's medical history. Patients recorded as "No" across 'Soft_drink,' 'HighBlood,' 'Overweight,' and 'Asthma' have a median of over 30 days of hospital stay. Low 'Complication_risk' and elective 'Initial_admin' also show the highest median compared to the other categories in each variable. These groups are interesting because they display the opposite of another assumption: medical conditions, high complication risk, and emergency admission should demonstrate more extended hospital stays, given these factors could make treatment more complex and thus lengthy.

C4. Data Transformation (Data Wrangling)

Data wrangling across some variables must ensure an accurate multiple linear regression model (Middleton, 2022a). I transformed some of the selected categorical variables, which are included in Figure 14 and Figure 15.

I re-expressed Boolean input variables with 'Yes' and 'No' to 1 and 0, respectively. Qualitative variables are split into separate variables and then one-hot encoded to consist of numeric values, also known as creating dummy variables or indicator variables. I dropped one column for each one-hot encoded variable to eliminate redundancy in the model (Lingsma et al., 2018). See Figure 14 below for re-expression of categorical variables. Figure 15 shows the code for one-hot encoding to selected categorical variables. To see these data-wrangling actions in detail, refer to the "Data Wrangling" section in the Jupyter Notebook file.

Figure 14

Re-expression of Categorical Variables Code and Output

```

# New column names variable
new_col_list = []

# Function for looping through target columns to re-express 'Yes/No' to 1/0
for col in reexp_cols:
    # Make new name and duplicate original column
    new_col = col +'_numeric'
    df[new_col] = df[col]

    # Make dictionary to change boolean to numeric
    new_dict = {new_col: {'No': 0, 'Yes': 1}}

    # Change all values in duplicated columns to ordinal encoding
    df.replace(new_dict, inplace=True)

    # Add to new_col_list
    new_col_list.append(new_col)

# Check that numeric columns were created by filtering and printing them
numeric_cols = [var for var in df.columns if 'numeric' in var]
print(df[numeric_cols].info())

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10000 entries, 0 to 9999
Data columns (total 8 columns):
 #   Column           Non-Null Count  Dtype  
 ---  -- 
 0   Soft_drink_numeric    10000 non-null   int64  
 1   HighBlood_numeric     10000 non-null   int64  
 2   Overweight_numeric    10000 non-null   int64  
 3   Arthritis_numeric     10000 non-null   int64  
 4   Diabetes_numeric      10000 non-null   int64  
 5   BackPain_numeric      10000 non-null   int64  
 6   Reflux_esophagitis_numeric 10000 non-null   int64  
 7   Asthma_numeric        10000 non-null   int64  
dtypes: int64(8)

# Show unique values in new_col_list
for col in new_col_list:
    print(col, df[col].unique())

```

Soft_drink_numeric [0 1]
 HighBlood_numeric [1 0]
 Overweight_numeric [0 1]
 Arthritis_numeric [1 0]
 Diabetes_numeric [1 0]
 BackPain_numeric [1 0]
 Reflux_esophagitis_numeric [0 1]
 Asthma_numeric [1 0]

Figure 15

One-Hot Encoding of Categorical Variables Code and Output

```

# Establish selection for dummy variable creation
target_dummy_vars = ['Initial_admin', 'Complication_risk', 'Services']
# Create dummy variables and drop one level to remove redundancy
dummies = pd.get_dummies(df[target_dummy_vars], prefix=target_dummy_vars, drop_first=True, dtype=int)
dummies.head()

Initial_admin_Emergency Admission Initial_admin_Observation Admission Complication_risk_Low Complication_risk_Medium Services_CT Scan Services_Intravenous Services_MRI
0 1 0 0 1 0 0 0
1 1 0 0 0 0 1 0
2 0 0 0 1 0 0 0
3 0 0 0 1 0 0 0
4 0 0 1 0 1 0 0

# Add dummies columns to df and make change permanent by reassigning df
df = pd.concat([df, dummies], axis=1)

# Create array for generated dummy variables and check that they made it to the dataframe
dummy_vars = ['Initial_admin_Emergency_Admission', 'Initial_admin_Observation_Admission', 'Complication_risk_Low', 'Complication_risk_Medium',
              'Services_CT_Scan', 'Services_Intravenous', 'Services_MRI']

print(df[dummy_vars].info())

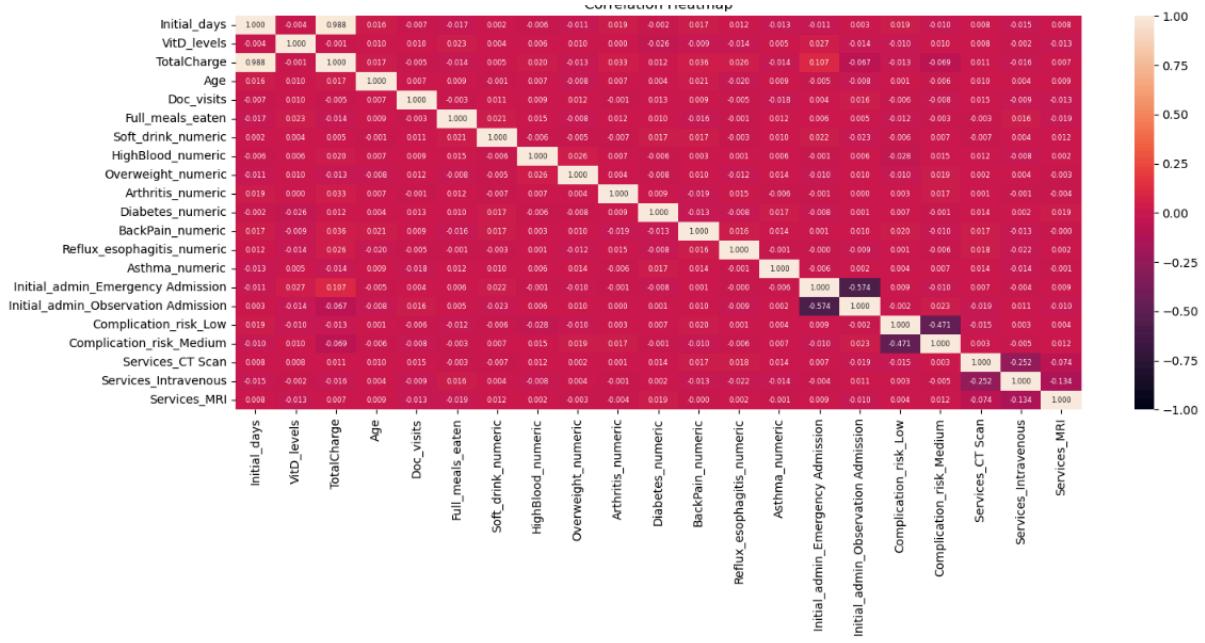
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10000 entries, 0 to 9999
Data columns (total 7 columns):
 #   Column           Non-Null Count  Dtype  
 --- 
 0   Initial_admin_Emergency_Admission 10000 non-null   int64 
 1   Initial_admin_Observation_Admission 10000 non-null   int64 
 2   Complication_risk_Low              10000 non-null   int64 
 3   Complication_risk_Medium          10000 non-null   int64 
 4   Services_CT_Scan                10000 non-null   int64 
 5   Services_Intravenous             10000 non-null   int64 
 6   Services_MRI                   10000 non-null   int64

```

After cleaning and wrangling the selected data, I created a correlation matrix of dependent and independent variables to check the linearity assumption. Figure 16 below shows the correlation matrix that indicates that many of the patient-level variables do not show strong linear correlations with 'Initial_days' except for 'TotalCharge' where the correlation coefficient is .99. For the sake of the exercise, we will move forward with these knowing these variables have violated the assumption of linearity.

Figure 16

Correlation Matrix



C5. CSV File of Prepared Data

Please refer to the generated CSV file titled 'Cleaned_Wrangle_MLR_Dataframe.'

Part IV: Model Comparison and Analysis

D1. Initial Model

Figure 17

Model A: Initial Multiple Linear Regression Model Code & Summary

```

# Establish the explanatory columns
expl_cols = df[['VitD_levels', 'TotalCharge', 'Age', 'Doc_visits', 'Full_meals_eaten', 'Complication_risk_Low',
                 'Complication_risk_Medium', 'Initial_admin_Emergency Admission', 'Initial_admin_Observation Admission',
                 'Services_CT Scan', 'Services_Intravenous', 'Services_MRI', 'Soft_drink_numeric', 'HighBlood_numeric',
                 'Overweight_numeric', 'Arthritis_numeric', 'Diabetes_numeric', 'BackPain_numeric', 'Reflux_esophagitis_numeric',
                 'Asthma_numeric']]]

y = df[['Initial_days']]
X = expl_cols.assign(const=1)

model_A = sm.OLS(y, X).fit()

print(model_A.summary())

```

OLS Regression Results

	coef	std err	t	P> t	[0.025	0.975]
VitD_levels	0.0024	0.004	0.599	0.549	-0.005	0.010
TotalCharge	0.0122	3.76e-06	3242.730	0.000	0.012	0.012
Age	-0.0004	0.000	-1.066	0.287	-0.001	0.000
Doc_visits	0.0135	0.008	1.746	0.081	-0.002	0.029
Full_meals_eaten	-0.0100	0.008	-1.238	0.216	-0.026	0.006
Complication_risk_Low	5.0380	0.023	223.842	0.000	4.994	5.082
Complication_risk_Medium	5.0346	0.019	271.520	0.000	4.998	5.071
Initial_admin_Emergency Admission	-6.2691	0.020	-315.380	0.000	-6.308	-6.230
Initial_admin_Observation Admission	0.0213	0.023	0.922	0.357	-0.024	0.067
Services_CT Scan	0.0028	0.026	0.108	0.914	-0.048	0.053
Services_Intravenous	0.0013	0.018	0.072	0.943	-0.035	0.037
Services_MRI	0.0292	0.043	0.678	0.498	-0.055	0.114
Soft_drink_numeric	-0.0246	0.019	-1.328	0.184	-0.061	0.012
HighBlood_numeric	-1.3766	0.016	-83.457	0.000	-1.409	-1.344
Overweight_numeric	0.0156	0.018	0.874	0.382	-0.019	0.051
Arthritis_numeric	-0.8867	0.017	-52.397	0.000	-0.920	-0.854
Diabetes_numeric	-0.9326	0.018	-51.270	0.000	-0.968	-0.897
BackPain_numeric	-1.0486	0.016	-63.570	0.000	-1.081	-1.016
Reflux_esophagitis_numeric	-0.7150	0.016	-43.416	0.000	-0.747	-0.683
Asthma_numeric	-0.0050	0.018	-0.279	0.781	-0.040	0.030
const	-28.7021	0.091	-315.271	0.000	-28.881	-28.524
Omnibus:	481.369	Durbin-Watson:	1.975			
Prob(Omnibus):	0.000	Jarque-Bera (JB):	383.476			
Skew:	-0.396	Prob(JB):	5.36e-84			
Kurtosis:	2.459	Cond. No.	6.48e+04			

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
[2] The condition number is large, 6.48e+04. This might indicate that there are strong multicollinearity or other numerical problems.

The model (Figure 17) using the Ordinary Least Squares method (OLS) shows an adjusted R-value of .999, which indicates that 99% of the variance of 'Initial_days' is explained by the combination of the variables. Two things stand out upon viewing the model summary. First, we see a multicollinearity error telling us we must check which variables are causing it. Multicollinearity among the explanatory variables affects the quality of the model. Second, we see that R^2 and adjusted R^2 are the same values. We set an alpha of 0.05 and see several of the variables' p-values are above the significance threshold.

Before reducing the model with backward stepwise elimination, we could check for multicollinearity using each variable's variance inflation factor (VIF). Removing these variables recursively by removing one at a time, starting with the highest VIF, will help us meet the non-multicollinearity assumption of a linear regression model. Addressing multicollinearity is essential because correlated predictor variables can cause unreliable coefficients and make models unstable (Larose & Larose, 2019). Figure 18 below shows the code and output for the variable inflation factors across all our explanatory variables. We set a VIF threshold of 10 to address severe multicollinearity, and anything over 5 and under 10 presents moderate collinearity. 'VitD_levels' and 'Doc_visits' have the highest VIFs. To see the recursion outputs, please refer to the "Addressing Multicollinearity" section in the Jupyter notebook titled "D208_MultiLinearRegression.ipynb." After recursively removing 'VitD_levels' and 'Doc_visits,' we have the VIFs in Figure 19.

Figure 18

First VIF Code and Output

```

: from statsmodels.stats.outliers_influence import variance_inflation_factor
vif_check = expl_cols
X = vif_check
# Create dataframe to load variable name and VIF scores
vif_df = pd.DataFrame()
# Create column called feature to list all variables
vif_df["feature"] = X.columns
# Creates column called 'VIF' to list all corresponding VIF scores
vif_df["VIF"] = [variance_inflation_factor(X.values, i)
for i in range(len(X.columns))]

print(vif_df)

```

	feature	VIF
0	VitD_levels	31.116574
1	TotalCharge	6.769467
2	Age	7.376655
3	Doc_visits	19.882159
4	Full_meals_eaten	1.983054
5	Complication_risk_Low	1.619661
6	Complication_risk_Medium	2.327733
7	Initial_admin_Emergency Admission	3.016369
8	Initial_admin_Observation Admission	1.952083
9	Services_CT Scan	1.233109
10	Services_Intravenous	1.585888
11	Services_MRI	1.072191
12	Soft_drink_numeric	1.347784
13	HighBlood_numeric	1.691125
14	Overweight_numeric	3.392016
15	Arthritis_numeric	1.556067
16	Diabetes_numeric	1.375266
17	BackPain_numeric	1.700139
18	Reflux_esophagitis_numeric	1.695831
19	Asthma_numeric	1.405584

Figure 19

VIFs After Removing 'VitD_levels' and 'Doc_visits'

```

: max_val = vif_df['VIF'].max()
max_name = vif_df['feature'][vif_df['VIF'] == max_val].values[0]
max_index = vif_df['feature'][vif_df['VIF'] == max_val].index[0]
print(max_val, max_name, max_index)

vif_df.drop(max_index, inplace=True)
print(vif_df)

19.882158972165122 Doc_visits 3
      feature      VIF
1      TotalCharge  6.769467
2          Age  7.376655
4      Full_meals_eaten  1.983054
5      Complication_risk_Low  1.619661
6      Complication_risk_Medium  2.327733
7  Initial_admin_Emergency Admission  3.016369
8  Initial_admin_Observation Admission  1.952083
9      Services_CT Scan  1.233109
10     Services_Intravenous  1.585808
11     Services_MRI  1.072191
12     Soft_drink_numeric  1.347784
13     HighBlood_numeric  1.691125
14     Overweight_numeric  3.392016
15     Arthritis_numeric  1.556067
16     Diabetes_numeric  1.375266
17     BackPain_numeric  1.700139
18     Reflux_esophagitis_numeric  1.695831
19     Asthma_numeric  1.405584

```

The output above shows moderate multicollinearity between 'TotalCharge' and 'Age.' Since the VIF threshold we set is 10 and these are below 10, we will keep them for now. Most notably, we know that 'TotalCharge' has a high linear correlation with 'Initial_days,' so keeping it in the model is useful.

D2. Linear Model Reduction Justification

After fitting the reduced set of variables with a new model, we get the summary in Figure 20.

Figure 20

Model After VIF Reduction

OLS Regression Results						
Dep. Variable:	Initial_days	R-squared:	0.999			
Model:	OLS	Adj. R-squared:	0.999			
Method:	Least Squares	F-statistic:	5.856e+05			
Date:	Tue, 06 Aug 2024	Prob (F-statistic):	0.00			
Time:	21:30:42	Log-Likelihood:	-12072.			
No. Observations:	10000	AIC:	2.418e+04			
Df Residuals:	9981	BIC:	2.432e+04			
Df Model:	18					
Covariance Type:	nonrobust					
	coef	std err	t	P> t	[0.025	0.975]
TotalCharge	0.0122	3.76e-06	3242.586	0.000	0.012	0.012
Age	-0.0004	0.000	-1.047	0.295	-0.001	0.000
Full_meals_eaten	-0.0099	0.008	-1.231	0.218	-0.026	0.006
Complication_risk_Low	5.0374	0.023	223.820	0.000	4.993	5.082
Complication_risk_Medium	5.0343	0.019	271.512	0.000	4.998	5.071
Initial_admin_Emergency Admission	-6.2682	0.020	-315.444	0.000	-6.307	-6.229
Initial_admin_Observation Admission	0.0222	0.023	0.962	0.336	-0.023	0.067
Services_CT Scan	0.0034	0.026	0.132	0.895	-0.047	0.054
Services_Intravenous	0.0010	0.018	0.056	0.955	-0.035	0.037
Services_MRI	0.0279	0.043	0.648	0.517	-0.057	0.112
Soft_drink_numeric	-0.0242	0.019	-1.305	0.192	-0.061	0.012
HighBlood_numeric	-1.3763	0.016	-83.438	0.000	-1.409	-1.344
Overweight_numeric	0.0161	0.018	0.901	0.368	-0.019	0.051
Arthritis_numeric	-0.8867	0.017	-52.394	0.000	-0.920	-0.854
Diabetes_numeric	-0.9324	0.018	-51.280	0.000	-0.968	-0.897
BackPain_numeric	-1.0485	0.016	-63.561	0.000	-1.081	-1.016
Reflux_esophagitis_numeric	-0.7153	0.016	-43.435	0.000	-0.748	-0.683
Asthma_numeric	-0.0055	0.018	-0.309	0.758	-0.041	0.030
const	-28.5919	0.041	-690.119	0.000	-28.673	-28.511
<hr/>						
Omnibus:	482.804	Durbin-Watson:	1.974			
Prob(Omnibus):	0.000	Jarque-Bera (JB):	383.918			
Skew:	-0.396	Prob(JB):	4.30e-84			
Kurtosis:	2.458	Cond. No.	3.21e+04			

Notes:
 [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
 [2] The condition number is large, 3.21e+04. This might indicate that there are strong multicollinearity or other numerical problems.

The figure above shows that we have reduced multicollinearity yet still get an error. Since 'TotalCharge' and 'Age' have VIFs of about 7, we can move on from the multicollinearity note and proceed with model reduction. 'Age' has a high p-value, so that we can remove it during the backward stepwise elimination process.

Since we have identified p-values higher than the established alpha of 0.05, we can perform a feature selection method. My choice of model reduction is to do backward stepwise elimination and use the p-value as a guide for the inclusion or exclusion of variables (Larose & Larose, 2019). Backward stepwise elimination removes one variable with a high p-value, re-runs the model, and identifies how the model improves or changes. Reducing variables in the model avoids overfitting with too many parameters that can cause noise with a different dataset instead of the training dataset (Larose & Larose, 2019). We can identify and select only the statistically relevant variables. You can see the recursive outputs of the

backward stepwise elimination process by referring to the "Backward Stepwise Elimination" section in the Jupyter Notebook titled "D208_MultiLinearRegression.ipynb."

D3. Reduced Model

Figure 21 below shows the reduced model after performing backward stepwise elimination. The Model B has eight variables plus the constant. The adjusted R² value stayed the same at .99.

Figure 21

Model B: After Backward Stepwise Elimination

OLS Regression Results						
Dep. Variable:	Initial_days	R-squared:	0.999			
Model:	OLS	Adj. R-squared:	0.999			
Method:	Least Squares	F-statistic:	1.171e+06			
Date:	Wed, 07 Aug 2024	Prob (F-statistic):	0.00			
Time:	17:19:21	Log-Likelihood:	-12075.			
No. Observations:	10000	AIC:	2.417e+04			
Df Residuals:	9990	BIC:	2.424e+04			
Df Model:	9					
Covariance Type:	nonrobust					
	coef	std err	t	P> t	[0.025	0.975]
TotalCharge	0.0122	3.76e-06	3244.808	0.000	0.012	0.012
Complication_risk_Low	5.0384	0.022	223.993	0.000	4.994	5.082
Complication_risk_Medium	5.0354	0.019	271.799	0.000	4.999	5.072
Initial_admin_Emergency Admission	-6.2797	0.016	-385.404	0.000	-6.312	-6.248
HighBlood_numeric	-1.3761	0.016	-83.486	0.000	-1.408	-1.344
Arthritis_numeric	-0.8869	0.017	-52.419	0.000	-0.920	-0.854
Diabetes_numeric	-0.9332	0.018	-51.363	0.000	-0.969	-0.898
BackPain_numeric	-1.0486	0.016	-63.635	0.000	-1.081	-1.016
Reflux_esophagitis_numeric	-0.7152	0.016	-43.461	0.000	-0.747	-0.683
const	-28.6076	0.029	-993.699	0.000	-28.664	-28.551
<hr/>						
Omnibus:	482.377	Durbin-Watson:	1.974			
Prob(Omnibus):	0.000	Jarque-Bera (JB):	384.450			
Skew:	-0.397	Prob(JB):	3.29e-84			
Kurtosis:	2.459	Cond. No.	2.35e+04			
<hr/>						
Notes:						
[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.						
[2] The condition number is large, 2.35e+04. This might indicate that there are strong multicollinearity or other numerical problems.						

The figure above shows a multicollinearity note, and rechecking the variance inflation factors shows that the constant has a high VIF (Figure 22). Since we must keep the Intercept ('const') in the model so that the fitted line is not biased, we ignore the multicollinearity note this time.

Figure 22

Final Multicollinearity Check

```

# Create dataframe to load variable name and VIF scores
vif_df = pd.DataFrame()

# Create column called feature to list all variables
vif_df["feature"] = X.columns
# Creates column called 'VIF' to list all corresponding VIF scores

vif_df["VIF"] = [variance_inflation_factor(X.values, i) for i in range(len(X.columns))]

# vif_df = vif_df.set_index('feature')
print(vif_df)

```

	feature	VIF
0	TotalCharge	1.023009
1	Complication_risk_Low	1.290980
2	Complication_risk_Medium	1.296056
3	Initial_admin_Emergency_Admission	1.011844
4	HighBlood_numeric	1.001273
5	Arthritis_numeric	1.002434
6	Diabetes_numeric	1.000648
7	BackPain_numeric	1.002556
8	Reflux_esophagitis_numeric	1.001217
9	const	12.637069

E1. Models Comparison

Figure 23

Model A (Initial) vs. Model B (Reduced)

OLS Regression Results											
Dep. Variable:	Initial_days	R-squared:	0.999	Model:	OLS	Adj. R-squared:	0.999	Method:	Least Squares	F-statistic:	5.271e+05
Date:	Thu, 01 Aug 2024	Prob (F-statistic):	0.00	Date:	Wed, 07 Aug 2024	Prob (F-statistic):	1.171e+06	Time:	17:19:21	Log-Likelihood:	-12875.
No. Observations:	10000	AIC:	2.418e+04	No. Observations:	10000	AIC:	2.417e+04	Df Residuals:	9990	BIC:	2.424e+04
Df Model:	2	Df Residuals:	9990	Df Model:	9	Df Residuals:	9990	Covariance Type:	nonrobust	Covariance Type:	nonrobust
Dep. Variable:	Initial_days	R-squared:	0.999	Model:	OLS	Adj. R-squared:	0.999	Method:	Least Squares	F-statistic:	1.171e+06
vitD_levels	0.0024	0.004	0.599	0.549	-0.005	0.010	0.010	Date:	Wed, 07 Aug 2024	Prob (F-statistic):	0.00
TotalCharge	0.0122	3.76e-06	3242.730	0.000	0.012	0.012	0.012	Time:	17:19:21	Log-Likelihood:	-12875.
Age	-0.0001	0.000	-1.066	0.287	-0.001	0.000	0.000	No. Observations:	10000	AIC:	2.417e+04
Doc_visits	0.0135	0.008	0.106	0.000	-0.002	0.009	0.009	Df Residuals:	9990	BIC:	2.424e+04
Full_meals_eaten	-0.0100	0.008	-1.238	0.216	-0.026	0.006	0.006	Df Model:	9	Covariance Type:	nonrobust
Complication_risk_Low	0.0380	0.023	223.842	0.000	4.994	5.082	5.082	Df Residuals:	9990	Covariance Type:	nonrobust
Complication_risk_Medium	0.0346	0.019	271.520	0.000	4.998	5.071	5.071	Df Model:	9	Covariance Type:	nonrobust
Initial_admin_Emergency_Admission	-0.2691	0.020	-315.388	0.000	-6.308	-6.230	-6.230	Df Residuals:	9990	Covariance Type:	nonrobust
Initial_admin_Observation_Admission	0.0213	0.023	0.922	0.357	-0.024	0.067	0.067	Df Model:	9	Covariance Type:	nonrobust
Services_CT_Scan	0.0028	0.026	0.108	0.914	-0.048	0.053	0.053	Df Residuals:	9990	Covariance Type:	nonrobust
Services_Intravenous	0.0013	0.018	0.072	0.943	-0.035	0.037	0.037	Df Model:	9	Covariance Type:	nonrobust
Services_MR	0.0292	0.043	0.678	0.498	-0.055	0.114	0.114	Df Residuals:	9990	Covariance Type:	nonrobust
Services_nu	0.0019	0.019	-1.024	0.261	-0.031	0.042	0.042	Df Model:	9	Covariance Type:	nonrobust
HighBlood_numeric	-0.3766	0.016	-83.457	0.000	-1.489	-1.344	-1.344	Df Residuals:	9990	Covariance Type:	nonrobust
Overweight_numeric	0.0156	0.018	0.874	0.382	-0.019	0.051	0.051	Df Model:	9	Covariance Type:	nonrobust
Arthritis_numeric	-0.8867	0.017	-52.397	0.000	-0.920	-0.854	-0.854	Df Residuals:	9990	Covariance Type:	nonrobust
Diabetes_numeric	-0.9326	0.018	-51.270	0.000	-0.968	-0.897	-0.897	Df Model:	9	Covariance Type:	nonrobust
BackPain_numeric	-1.0486	0.016	-63.570	0.000	-1.081	-1.016	-1.016	Df Residuals:	9990	Covariance Type:	nonrobust
Reflux_esophagitis_numeric	-0.7150	0.016	-43.416	0.000	-0.747	-0.683	-0.683	Df Model:	9	Covariance Type:	nonrobust
Asthma_numeric	-0.0050	0.018	-0.279	0.781	-0.048	0.030	0.030	Df Residuals:	9990	Covariance Type:	nonrobust
const	-28.7021	0.091	-315.271	0.000	-28.881	-28.524	-28.524	Df Model:	9	Covariance Type:	nonrobust
<hr/>											
Omnibus:	481.369	Durbin-Watson:	1.975		Omnibus:	482.377	Durbin-Watson:	1.974			
Prob(Omnibus):	0.000	Jarque-Bera (JB):	383.476		Prob(Omnibus):	0.000	Jarque-Bera (JB):	384.450			
Skew:	-0.395	Prob(JB):	5.36e-84		Skew:	-0.397	Prob(JB):	3.29e-84			
Kurtosis:	2.459	Cond. No.	6.48e+04		Kurtosis:	2.459	Cond. No.	2.35e+04			

Notes:
 [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
 [2] The condition number is large, 6.48e+04. This might indicate that there are strong multicollinearity or other numerical problems.

Notes:
 [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
 [2] The condition number is large, 2.35e+04. This might indicate that there are strong multicollinearity or other numerical problems.

The adjusted R² values (Figure 23) remain unchanged at .99, meaning that the predictor variables can explain 99% of the variation in 'Initial_days.' The F-statistic probability is 0 on both models, which also tells us that overall, the variables as a group in each regression model are significant. The AIC and BIC differences between the two models are extremely minimal, with Model B being the 'better' model with slightly fewer values. The residual standard error and residuals are also very similar in value

between the two models, as seen in Figure 24. A closer look at the numbers shows Model B edging out Model A for better RSE at the hundredth thousandth place, yet this insight is negligible.

The main differences in the model summaries lie within the individual variable's statistical significance in combination. In the initial model, there are variables with p-values well above the alpha, adding to the noise and risking overfitting the model. The final model includes only individual statistically significant variables, making it more reliable. Because Model B consists of only statistically significant variables compared to Model A, we will move forward with the reduced model (Model B) because of the significance of its explanatory variables.

Let us revisit the four assumptions of a multiple linear regression model mentioned in section B1. So far, we know neither model passes the assumption of linearity due to the overview of the correlation matrix generation in section C4. Only one explanatory variable shows a linear correlation to the predictor variable. In contrast, the rest of the variables that were added for the exercise denote a very low or nearly no linear relationship to the response variables. Both Model A and Model B violate the linearity assumption.

There is a difference in the multicollinearity assumption between Model A and Model B. Model A exhibited severe multicollinearity compared to Model B. Part of the reduction method for Model B was to remove variables that exhibited high VIFs. Thus, the reduced model consists of variables with VIFs less than 5 to prevent severe multicollinearity, which meets the second assumption of this exercise.

The third assumption is homoscedasticity, plotted for Model B in Figure 25. The residual plot shows horizontal bands evenly distributed across the graph. Although the plot does not exhibit a random pattern, the variance between 0 and the points does not increase as it does with heteroscedasticity (*Qualtrics*, n.d.). The placement of points is constant around 0 between 1 and -2, making the variance predictable. Model B meets the third assumption of homoscedasticity.

The QQ plot of predicted values against Model B's residuals is shown in Figure 26. It does not demonstrate normality because quite a few points do not fall on the 45-degree line. These points are very close and appear horizontal, so I checked the top ten largest and smallest residual values to make sure that the values are continuous. The calculation in Figure 26 showed that they are continuous but incremented very slightly from one another. The generated histogram in Figure 27 confirms that the residuals are not normally distributed, so Model B does not meet this normality assumption.

E2. Screenshots of Outputs and Calculation

Figure 24

Residual Standard Error Calculations (Model A and Model B)

```
# Calculate Residual Standard Error
model_A_rse = np.sqrt(model_A.mse_resid)
model_B_rse = np.sqrt(model_B.mse_resid)

print('Model A RSE: {}'.format(model_A_rse))
print('Model B RSE: {}'.format(model_B_rse))

Model A RSE: 0.8098814229317172
Model B RSE: 0.8098483722308781
```

Figure 25

Model B Residual Plot

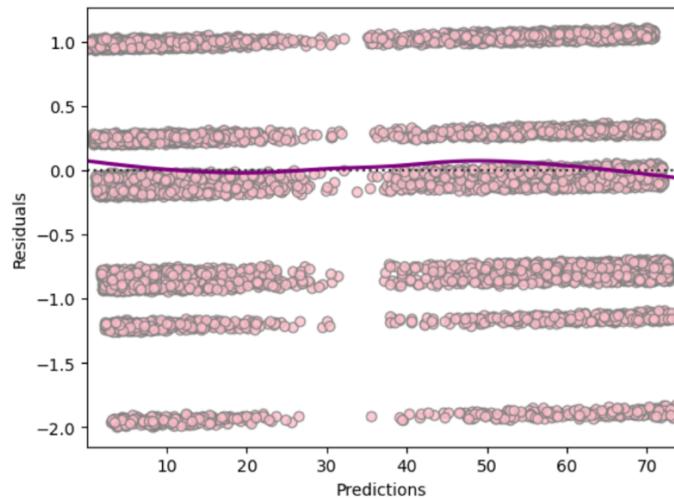


Figure 26

Model B QQ Plot & Residual Values

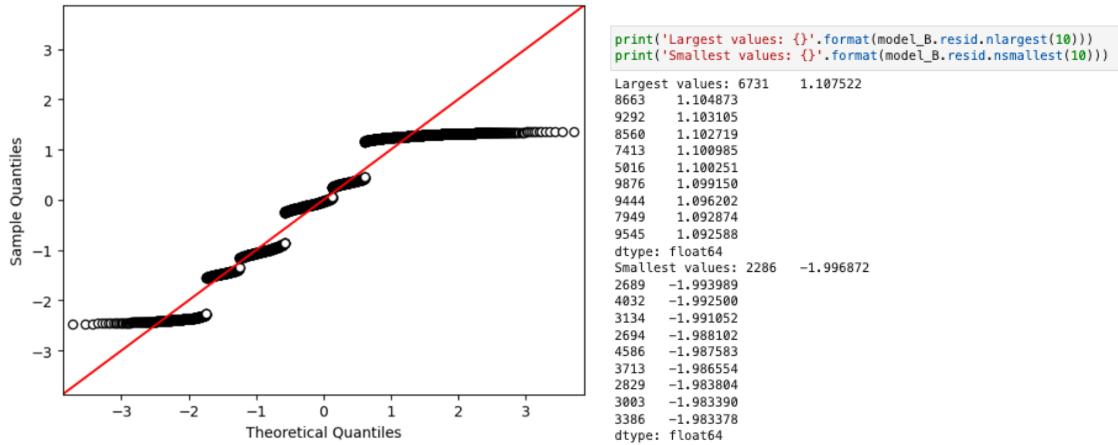
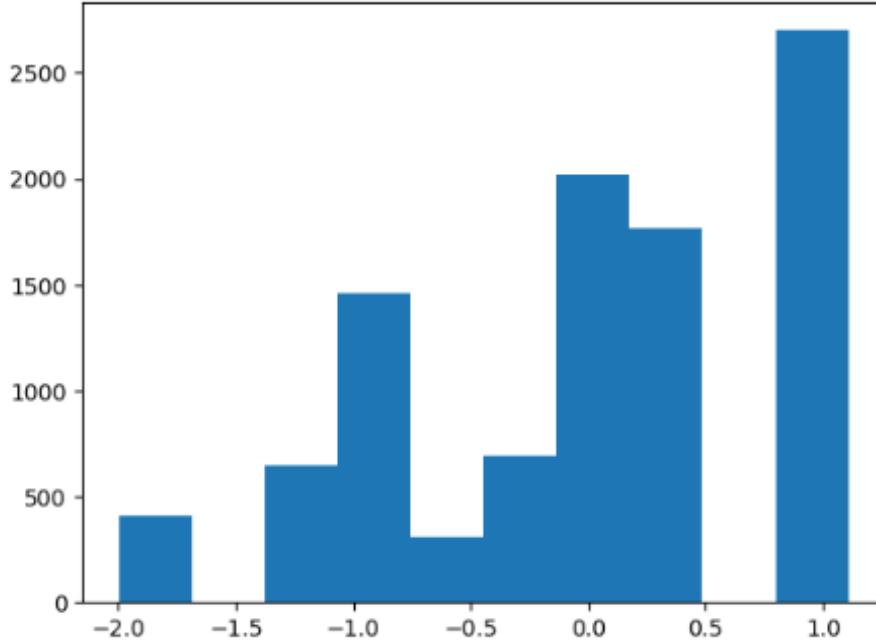


Figure 27

Model B Residual Histogram Plot



E3. Executable Code

See the file titled "Multiple_Linear_Regression_Notebook.ipynb" to see all the code generated for this paper.

Part V: Data Summary and Implications

F1. Results of Data Analysis

Figure 28

Regression Equation

$$\text{Hospitalization Days} = -28.6076 + 0.0122 \text{ (Average Daily Charge)} + 5.0384 \text{ (Low Complication Risk)} + 5.0354 \text{ (Medium Complication Risk)} - 6.2797 \text{ (Emergency Admission)} - 1.3761 \text{ (High Blood)} - 0.8869 \text{ (Arthritis)} - 0.9332 \text{ (Diabetes)} - 1.0486 \text{ (Back Pain)}$$

The equation above (Figure 28) reflects Model B. The naming convention has been changed slightly for improved readability. 'Hospitalization Days' is the response variable 'Initial_days' or y that equals the following expression. The regression coefficient interpretations are as follows:

- Holding all other variables constant, one unit *increase in average daily charge* is associated with 0.0122 days *increase* in a patient's hospitalization stay.
- Holding all other variables constant, patients who have low complication risk, on average, stay at the hospital 5.0384 *more* days.
- Holding all other variables constant, patients who have medium complication risk, on average, stay at the hospital 5.0354 *more* days.
- Holding all other variables constant, patients who have emergency admission status, on average, stay at the hospital 6.2797 *fewer* days.
- Holding all other variables constant, patients who have high blood pressure, on average, stay at the hospital 1.3761 *fewer* days than patients who don't.
- Holding all other variables constant, patients who have arthritis stay at the hospital 0.8869 *fewer* days than patients who don't.
- Holding all other variables constant, patients who have diabetes stay at the hospital 0.9332 *fewer* days than patients who don't.
- Holding all other variables constant, patients who have back pain stay at the hospital 1.0486 *fewer* days than patients who don't.

This reduced model and combined explanatory variables demonstrate a vital statistical significance in helping us find some patient factors for the research question. Model B's low F-statistic probability and meaningful p-values signal statistical significance within the overall model.

However, I would not call Model B practically significant for the reasons below. Some are reiterations of insights taken from other sections in this paper.

The first reason Model B is not practically significant is because Model B does not meet all four assumptions of linear regression. To echo section E1, Model B meets the multicollinearity and homoscedastic assumptions but fails to meet linearity and normality assumptions. Issues with homoscedasticity and multivariate normality, as mentioned in section B1, make models unreliable. Not meeting all the conditions lowers the confidence that this model is helpful in a practical sense (Middleton, 2022b).

The second reason is that when it comes to the individual explanatory variable significance, we know that 'TotalCharge' is the only variable with a linear correlation to the response variable 'Initial_days.' This is not a helpful variable when providing insights about patients' demographic information despite being counted as a patient-hospital interaction factor. Clearly, 'TotalCharge' and 'Initial_days' would be positively correlated since more extended hospital stays mean a patient would have to pay for more days. It is only helpful as a significant variable because it provides a correlation weight to the model and lends significance to the other explanatory variables. If we were to run the model without 'TotalCharge,' like in Figure 29, the model would be rendered useless with an extremely low adjusted R^2 and the remaining variables' significance values inflate above the target alpha.

Figure 29

Model Without 'TotalCharge'

OLS Regression Results						
Dep. Variable:	Initial_days	R-squared:	0.001			
Model:	OLS	Adj. R-squared:	0.001			
Method:	Least Squares	F-statistic:	1.667			
Date:	Fri, 09 Aug 2024	Prob (F-statistic):	0.101			
Time:	21:54:00	Log-Likelihood:	-46881.			
No. Observations:	10000	AIC:	9.378e+04			
Df Residuals:	9991	BIC:	9.385e+04			
Df Model:	8					
Covariance Type:	nonrobust					
	coef	std err	t	P> t	[0.025	0.975]
Complication_risk_Low	1.1295	0.730	1.548	0.122	-0.301	2.559
Complication_risk_Medium	-0.1144	0.599	-0.191	0.849	-1.290	1.061
Initial_admin_Emergency Admission	-0.6082	0.526	-1.156	0.248	-1.639	0.423
HighBlood_numeric	-0.3236	0.535	-0.605	0.545	-1.373	0.726
Arthritis_numeric	1.0488	0.549	1.910	0.056	-0.028	2.125
Diabetes_numeric	-0.1484	0.590	-0.251	0.801	-1.305	1.008
BackPain_numeric	0.9163	0.535	1.713	0.087	-0.132	1.965
Reflux_esophagitis_numeric	0.6218	0.534	1.164	0.245	-0.425	1.669
const	33.7388	0.696	48.455	0.000	32.374	35.104
Omnibus:	41272.257	Durbin-Watson:		0.161		
Prob(Omnibus):	0.000	Jarque-Bera (JB):		1284.428		
Skew:	0.070	Prob(JB):		1.23e-279		
Kurtosis:	1.250	Cond. No.		5.61		

Alternatively, doing a simple linear regression with just 'TotalCharge' and 'Initial_days' would meet all four assumptions. However, to iterate, 'TotalCharge' does not help answer the entire research question because it is not a targeted factor of the patient medical or demographic background, we deem insightful.

There are several limitations to consider in this linear regression exercise. The overall dataset did not provide multiple variables that showed high correlation coefficients to the response variable. The model immediately could not meet the linearity assumption for linear regression, lowering the confidence in its usefulness. Regarding data cleaning and wrangling, I decidedly kept outliers, which all seemed realistic across all the variables. By including the outliers, I may have affected coefficients and potentially distorted residuals (Middleton, 2022b). Overall, the general dataset does not provide in-depth information about each variable. For example, a few categories are listed for 'Services.' 'TotalCharge,' while it increases as 'Initial_days' increase, could vary per patient's hospitalized treatments or insurance copay. Some vague dictionary definitions limit the dataset knowledge and consideration across variables. I also limited myself to backward stepwise elimination and VIF observations regarding the model reduction process. Some consider stepwise elimination faulty because individual variable selection may not necessarily help the model's overall significance. Smith (2018) gives an example of forward stepwise

elimination, where choosing a fifth explanatory variable based on the four selected variables' p-values does not tell us that these five variables will result in the highest R^2 . However, other methods are meticulous and time-consuming, especially when there are many variables to examine.

F2. Actionable Recommendations

Our research question is: *What factors regarding a patient's hospital activities, medical history, and demographic background correlate with the length of their initial hospital stay?* Model B's statistical significance shows us that some demographic factors can influence the length of a patient's hospital stay. However, since it does not meet the linear regression assumptions, the usefulness of Model B ends here. Although we discovered some factors correlated with a hospital stay, we should also be able to use the model confidently with data that is not the training data. Patient demographic information changes and shifts, so moving forward with a definitive model will also provide a solid foundation that can be observed over time.

The model is not practically significant, and therefore, subsequent recommendations address the limitations mentioned in section F1. We must reconsider the dataset and expand on the variables to build a model that can assertively predict the length of hospital stay based on patient demographic information. For example, we need to find other variables with a linear correlation to 'Initial_days.' We should also consider the peculiarities of the bivariate statistics mentioned in Section C3 and how they might affect the regression model. 'Initial_days' is bimodal, with two different data sets dispersed on either end of its histogram. We should figure out why there seem to be two general groups of patients. A deep dive into the definitions of some variables is also important and we should ask questions about limitations on some of them. Such questions could include the following: Why is 'Services' limited to only four types? How come most patients seem only to eat one meal? How are Vitamin D levels relevant to this dataset, and what are other vitamin types?

Finally, given what we know about the correlation between the explanatory variables and the response variable, we can explore other models. For example, since the variables violate the linearity assumption of a linear regression model, we know that this model is not an appropriate one to apply to our research question. An excellent start would be considering nonlinear models where the variables do not violate the linearity assumption.

Part VI: Demonstration

G. Sources for Third-Party Code

1. *Assessing model fit - Visualizing model fit.* (n.d.). [Video]. DataCamp.
<https://campus.datacamp.com/courses/introduction-to-regression-with-statsmodels-in-python/assessing-model-fit-e78fd9fe-6303-4048-8748-33b19c4222fe?ex=4>
2. *Pandas get_dummies()*. (n.d.). https://www.programiz.com/python-programming/pandas/methods/get_dummies
3. *Plotting many pie charts using a loop to create a single figure using matplotlib.* (n.d.). Stack Overflow. <https://stackoverflow.com/questions/75176180/plotting-many-pie-charts-using-a-loop-to-create-a-single-figure-using-matplotlib>
4. *Predicting again | Python.* (n.d.). DataCamp.
<https://campus.datacamp.com/courses/intermediate-regression-with-statsmodels-in-python/multiple-linear-regression-3?ex=10>
5. *seaborn.residplot — seaborn 0.13.2 documentation.* (n.d.).
<https://seaborn.pydata.org/generated/seaborn.residplot.html>
6. Sewell, W. (n.d.). *D208 Predictive Modeling Episode 1* [Video]. Panopto. Retrieved July 12, 2024, from

<https://wgu.hosted.panopto.com/Panopto/Pages/Auth/Login.aspx?instance=WGU&Auth=SessionView&panoptoState=276ead8c-2e90-419f-aab6-b1cb001544e3>

H. References

1. Bobbitt, Z. (2021, November 16). *The five assumptions of multiple linear regression*. Statology. <https://www.statology.org/multiple-linear-regression-assumptions/>
2. Bruce, P., Bruce, A., & Gedeck, P. (2020). *Practical Statistics for Data Scientists: 50+ essential concepts using R and Python*. <https://openlibrary.telkomuniversity.ac.id/home/catalog/id/163886/slug/practical-statistics-for-data-scientists-50-essential-concepts-using-r-and-python-2-e-.html>
3. Larose, C. D., & Larose, D. T. (2019). *Data science using Python and R*. John Wiley & Sons.
4. Lingsma, H. F., Bottle, A., Middleton, S., Kievit, J., Steyerberg, E. W., & De Mheen, P. J. M. (2018). Evaluation of hospital outcomes: the relation between length-of-stay, readmission, and mortality in a large international administrative database. *BMC Health Services Research*, 18(1). <https://doi.org/10.1186/s12913-018-2916-1>
5. Middleton, K. (2022a, November). *D208 - Webinar: Getting Started with D208 Part I (November 2022)* [Video]. Panopto. Retrieved July 13, 2024, from <https://wgu.hosted.panopto.com/Panopto/Pages/Viewer.aspx?id=15e09c73-c5aa-439d-852f-af47001b8970>
6. Middleton, K. (2022b, November). *D208 - Webinar: Getting Started with D208 Part II (November)* [Video]. Panopto. Retrieved July 28, 2024, from <https://wgu.hosted.panopto.com/Panopto/Pages/Viewer.aspx?id=39bbe2db-de7d-4bf5-913>

7. *Qualtrics*. (n.d.). <https://www.qualtrics.com/support/stats-iq/analyses/regression-guides/interpreting-residual-plots-improve-regression/>
8. Smith, G. (2018). Step away from stepwise. *Journal of Big Data*, 5(1).
<https://doi.org/10.1186/s40537-018-0143-6>
9. Tipton, K., Leas, B. F., Mull, N. K., Siddique, S. M., Greysen, S. R., Lane-Fall, M. B., & Tsou, A. Y. (2021, September 1). *Introduction*. Interventions to Decrease Hospital Length of Stay - NCBI Bookshelf. <https://www.ncbi.nlm.nih.gov/books/NBK574438/>